

STAT 621 Lecture Notes

Classification Examples

We've already seen several methods for classifying responses. Next we will work through an example or two to compare and contrast these procedures.

Example: Orange Juice Preference

The OJ data set in the ISLR package contains records of 1070 customer purchases of either Citrus Hill (CH) or Minute Maid (MM) orange juice. Seventeen other characteristics of the customer and product were also recorded. We'll reduce these a bit for this example and focus on predicting the brand purchased **Purchase** as a function of the following features.

WeekofPurchase: week of purchase

StoreID: store ID

PriceCH: price of Citrus Hill

PriceMM: price of Minute Maid

DiscCH: amount of discount for Citrus Hill

DiscMM: amount of discount for Minute Maid

SpecialCH: indicator for special on Citrus Hill (1=yes, 0=no)

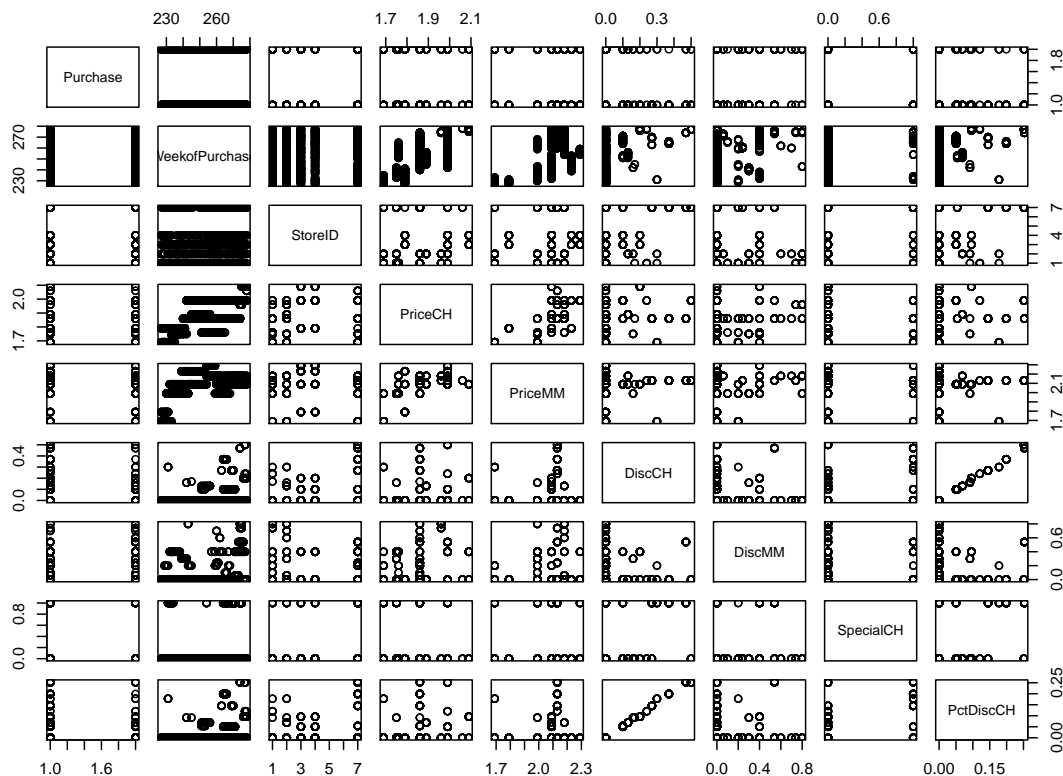
PctDiscCH: percentage of discount for Citrus Hill

Here's a bit of the data and a pairwise scatterplot.

```
# =====
library(MASS)
library(ISLR)
library(e1071)
library(caret)
library(klaR)

oj=OJ[,c(1:8,16)] # reduce the no. X's for this example
head(oj)
pairs(oj)
```

	Purchase	WeekofPurchase	StoreID	PriceCH	PriceMM	DiscCH	DiscMM	SpecialCH	PctDiscCH
1	CH	237	1	1.75	1.99	0.00	0.0	0	0.000000
2	CH	239	1	1.75	1.99	0.00	0.3	0	0.000000
3	CH	245	1	1.86	2.09	0.17	0.0	0	0.091398
4	MM	227	1	1.69	1.69	0.00	0.0	0	0.000000
5	CH	228	7	1.69	1.69	0.00	0.0	0	0.000000
6	CH	230	7	1.69	1.99	0.00	0.0	0	0.000000



We will consider predicting `Purchase` using the classification methods we've discussed recently. I'll split the data into a training set and test set. We will use about 75% of the data for training, and predict the response on the remaining 25%.

```
train = sample(1:nrow(oj), 800) # 800 = 75% of data for train
oj.train = oj[train, ]
oj.test = oj[-train, ]
```

Logistic Regression: First recall how this works. What is the model, and how do we use it to predict the response?

The logistic model is fit below. I use the `predict.glm` function to report the estimated probability of purchasing Minute Maid, $P(Y = 1)$, on the testing data. An estimated probability of greater than 0.5 is predicted as a Minute Maid purchase, otherwise it is predicted as a Citrus Hill purchase. The `confusionMatrix` function in the `caret` package is handy for summarizing the performance of the classifier.

```
> logistic1=glm(Purchase~., data=oj.train, family=binomial)
> summary(logistic1)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	4.78071	1.59829	2.991	0.002779	**
WeekofPurchase	-0.02873	0.00825	-3.483	0.000496	***
StoreID	-0.20337	0.03824	-5.318	1.05e-07	***
PriceCH	3.53055	1.19794	2.947	0.003207	**
PriceMM	-1.88928	0.75344	-2.508	0.012157	*
DiscCH	9.05908	17.19986	0.527	0.598405	
DiscMM	2.21239	0.41658	5.311	1.09e-07	***
SpecialCH	-0.23796	0.29758	-0.800	0.423902	
PctDiscCH	-19.94989	32.57272	-0.612	0.540225	

```
---
Null deviance: 1076.00 on 799 degrees of freedom
Residual deviance: 960.45 on 791 degrees of freedom
AIC: 978.45
```

```
> a=predict.glm(logistic1, oj.test, type="response")
> logistic.pred=factor(as.numeric(a>0.5))
> confusionMatrix(data = logistic.pred, reference = oj.test$Purchase)
```

Confusion Matrix and Statistics

	Reference	
Prediction	CH	MM
CH	142	60
MM	30	38

```
Accuracy : 0.6667
95% CI : (0.607, 0.7226)
No Information Rate : 0.637
P-Value [Acc > NIR] : 0.171400
```

```
Kappa : 0.2284
```

```
McNemar's Test P-Value : 0.002237 # test of table symmetry
```

```
Sensitivity : 0.8256 # prob predict CH given CH
Specificity : 0.3878 # prob predict MM given MM
Pos Pred Value : 0.7030
Neg Pred Value : 0.5588
Prevalence : 0.6370
Detection Rate : 0.5259
Detection Prevalence : 0.7481
Balanced Accuracy : 0.6067 # ave sens and spec
```

```
'Positive' Class : CH
```

Naive Bayes Model: Let's see how well this one does to predict **Purchase**. First a little review:

```
> nb1=NaiveBayes(Purchase~., data=oj.train, usekernel=T)
> nb.pred=predict(nb1, oj.test)
> confusionMatrix(data=nb.pred$class, reference=oj.test$Purchase)
```

Confusion Matrix and Statistics

	Reference	
Prediction	CH	MM
CH	162	83
MM	10	15

Accuracy : 0.6556
95% CI : (0.5956, 0.7121)
No Information Rate : 0.637
P-Value [Acc > NIR] : 0.2859

Kappa : 0.113

Mcnemar's Test P-Value : 8.264e-14

Sensitivity : 0.9419
Specificity : 0.1531
Pos Pred Value : 0.6612
Neg Pred Value : 0.6000
Prevalence : 0.6370
Detection Rate : 0.6000
Detection Prevalence : 0.9074
Balanced Accuracy : 0.5475

'Positive' Class : CH

Linear Discriminant Analysis: Recall how this works.

Here is the fit and some summaries. Since the response is binary, there's only one discriminant function (Best I can figure, R uses a scaling or normalizing to express the discriminator in terms of $p - 1$ functions).

```
> lda1=lda(Purchase~., data=oj.train)
> lda1
```

Prior probabilities of groups:


	CH	MM
	0.60125	0.39875

Group means:


	WeekofPurchase	StoreID	PriceCH	PriceMM	DiscCH	DiscMM	SpecialCH	PctDiscCH
CH	256.2599	4.413721	1.873306	2.100353	0.06688150	0.09692308	0.17671518	0.03523688
MM	251.9530	3.238245	1.864796	2.056301	0.02714734	0.17203762	0.09090909	0.01408230

Coefficients of linear discriminants:

	LD1
WeekofPurchase	-0.03415863
StoreID	-0.25148057
PriceCH	4.18384055
PriceMM	-2.55493185
DiscCH	9.65383732
DiscMM	2.75079142
SpecialCH	-0.21736277
PctDiscCH	-20.49016136



```
> plot(lda1)
> lda.pred=predict(object=lda1, newdata=oj.test)
> confusionMatrix(lda.pred$class, oj.test$Purchase)
```



Confusion Matrix and Statistics

	Reference	
Prediction	CH	MM
CH	147	63
MM	25	35

Accuracy : 0.6741
95% CI : (0.6146, 0.7296)
Kappa : 0.2311
McNemar's Test P-Value : 8.006e-05

Sensitivity : 0.8547
Specificity : 0.3571
Balanced Accuracy : 0.6059

Support Vector Machine: Finally we'll try making predictions using the SVM. This should be pretty fresh so I'll skip the review. Below I fit the SVM using the radial kernel. The parameters for the model, `gamma` and `cost` are estimated by a grid search for the combination with the lowest crossvalidated prediction error. The `tune` function in the `e1070` package makes this easy.

```
> tune.out=tune(svm, Purchase~., data=oj.train, kernel="radial", ranges=list(cost=seq(20,70,10), gamma=seq(.1,.5,.1)))
> summary(tune.out)
```

Parameter tuning of svm:

- sampling method: 10-fold cross validation

- best parameters:

```
cost gamma
 60    0.2
```

- best performance: 0.29875

- Detailed performance results:

	cost	gamma	error	dispersion	
1	50	0.1	0.31625	0.03537988	
2	60	0.1	0.31250	0.03486083	
3	70	0.1	0.30750	0.04216370	
4	80	0.1	0.30750	0.04377975	
5	90	0.1	0.30875	0.04332131	
6	50	0.2	0.30000	0.04409586	
7	60	0.2	0.29875	0.04348132	<-----
8	70	0.2	0.30125	0.04505013	
9	80	0.2	0.30000	0.04370037	
10	90	0.2	0.29875	0.04466309	
11	50	0.3	0.30250	0.04993051	
...	ETC				

```
> svm1=svm(Purchase~., data=oj.train, kernel="radial", gamma=.2, cost=50)
> summary(svm1)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: radial
cost: 50
```

```
Number of Support Vectors: 526
( 259 267 )
```

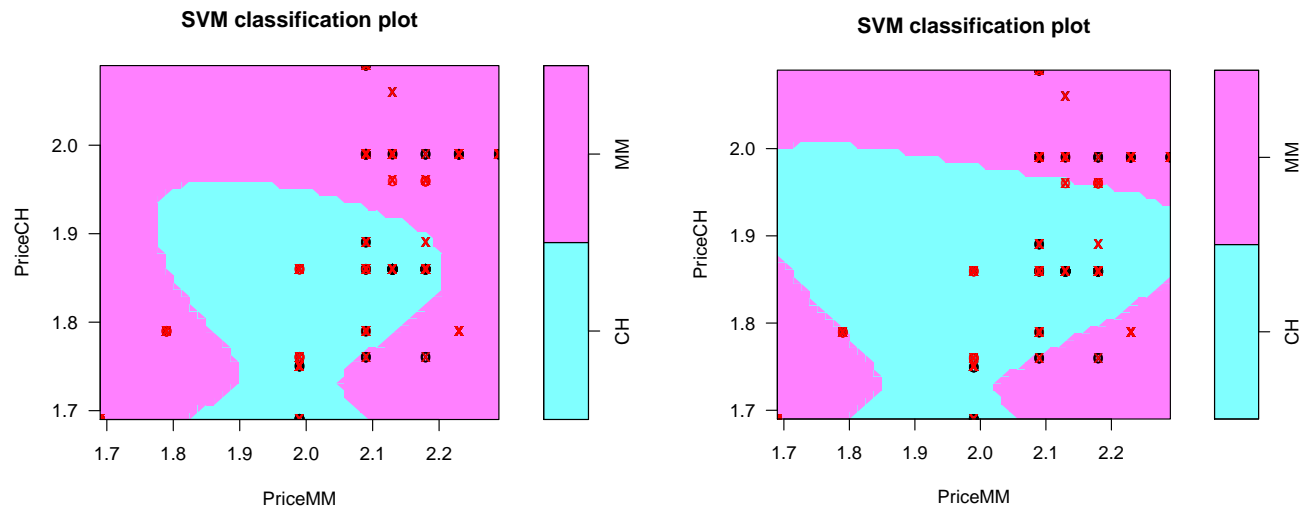
Number of Classes: 2

Levels:

```
CH MM
```

With $p > 2$, the `plot` function will show projections of the decision boundary projected onto the plane spanned by any two features. This can be done for specific values of the other features, basically showing a slice of the projection.

```
> plot(svm1,oj.train, PriceCH~PriceMM, slice=list(WeekofPurchase=250, StoreID=3))
> plot(svm1,oj.train, PriceCH~PriceMM, slice=list(WeekofPurchase=250, StoreID=5))
```



Now let's see how it does at classifying OJ purchases.

```
> svm.pred=predict(svm1,oj.test)
> confusionMatrix(svm.pred, oj.test$Purchase)
```

Confusion Matrix and Statistics

	Reference	
Prediction	CH	MM
CH	138	41
MM	34	57

Accuracy : 0.7222
 95% CI : (0.6647, 0.7748)
 Kappa : 0.3899
 McNemar's Test P-Value : 0.488422

Sensitivity : 0.8023
 Specificity : 0.5816
 ...
 Balanced Accuracy : 0.6920

Summarize the results.

All the methods were fairly similar in terms of prediction accuracy. But remember these results were based on a single split of the data into a training set and a test set. Maybe a better way to select a procedure would be to repeat this and choose the procedure that has the lowest error rate. In other words, let's crossvalidate!

```
K=5
cv.error = matrix(0,nrow=4,ncol=K)
set.seed(135)
folds = sample(1:K,nrow(oj),replace=T) # will give more or less same n each fold

for(i in 1:K)
{
  CV.train = oj[folds != i,]
  CV.test = oj[folds == i,]
  # logistic
  logistic.fit = glm(Purchase ~., data = CV.train, family = binomial)
  logistic.probs = predict(logistic.fit, CV.test, type = "response")
  logistic.pred = factor(as.numeric(logistic.probs>0.5))
  cv.error[1,i]=mean(logistic.pred != as.numeric(CV.test$Purchase)-1)
  # naive bayes
  nb.fit=NaiveBayes(Purchase~., data=CV.train, usekernel=T)
  nb.pred=predict(nb.fit, CV.test)
  cv.error[2,i]=mean(nb.pred$class != CV.test$Purchase)
  # LDA
  lda.fit=lda(Purchase~., data=CV.train)
  lda.pred=predict(lda.fit, CV.test)
  cv.error[3,i]=mean(lda.pred$class != CV.test$Purchase)
  # SVM -- use the parameter values found earlier
  svm.fit=svm(Purchase~., data=CV.train, kernel="radial", gamma=.2, cost=60)
  svm.pred=predict(svm.fit,CV.test)
  cv.error[4,i]=mean(svm.pred != CV.test$Purchase)
}

> row.names(cv.error)=c("logistic","NBayes","LDA","SVM")
> apply(cv.error, 1, mean)

logistic    NBayes      LDA      SVM
0.3532973 0.3757455 0.3479311 0.3003532
```

A pretty clear choice.